

Supplementary information for

Accelerating physics-informed neural networks for full waveform inversion using a hybrid quantum-classical finite-basis architecture

Hoang Anh Nguyen^{a,*}, Divakar Vashisth^b, and Ali Tura^a

^aDepartment of Geophysics, Colorado School of Mines, Golden, CO 80401, USA

^bDepartment of Energy Science and Engineering, Stanford University, Stanford, CA 94305, USA

*Corresponding author: hoanganh.nguyen@mines.edu

25th May 2026

Supplementary Note: Quantum background

For completeness, we summarize the basic quantum-computing concepts used by the hybrid architecture. A single qubit is a two-level quantum system with state

$$|\psi\rangle = \alpha|0\rangle + \beta|1\rangle, \quad |\alpha|^2 + |\beta|^2 = 1, \quad (\text{S1})$$

where $\alpha, \beta \in \mathbb{C}$ are probability amplitudes. An n -qubit register is represented by a normalized statevector in $\mathbb{C}^{2^n}$,

$$|\psi\rangle = \sum_{k=0}^{2^n-1} \alpha_k |k\rangle, \quad \sum_{k=0}^{2^n-1} |\alpha_k|^2 = 1. \quad (\text{S2})$$

The parameterized single-qubit rotations used in this work are

$$R_k(\theta) = e^{-i\theta \sigma_k/2}, \quad k \in \{x, y, z\}, \quad (\text{S3})$$

with explicit matrices

$$R_x(\theta) = \begin{pmatrix} \cos \frac{\theta}{2} & -i \sin \frac{\theta}{2} \\ -i \sin \frac{\theta}{2} & \cos \frac{\theta}{2} \end{pmatrix}, \quad R_y(\theta) = \begin{pmatrix} \cos \frac{\theta}{2} & -\sin \frac{\theta}{2} \\ \sin \frac{\theta}{2} & \cos \frac{\theta}{2} \end{pmatrix}, \quad R_z(\theta) = \begin{pmatrix} e^{-i\theta/2} & 0 \\ 0 & e^{i\theta/2} \end{pmatrix}. \quad (\text{S4})$$

The composition $R_z(\theta_3)R_y(\theta_2)R_x(\theta_1)$ spans $\text{SU}(2)$ and can prepare any single-qubit pure state up to a physically irrelevant global phase.

Entanglement between neighbouring qubits is introduced by the controlled-NOT (CNOT) gate,

$$\text{CNOT} = \begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 \end{pmatrix}, \quad (\text{S5})$$

which flips the target qubit if and only if the control qubit is in state $|1\rangle$.

Supplementary Figures




Figure S1: Comparison of parameter gradients $\partial f / \partial \theta_q$ computed by our JAX statevector implementation and PennyLane's `default.qubit` backend across eight circuit configurations (2–8 qubits, 1–6 layers). All values fall on the identity line $y = x$ with maximum absolute deviation $\Delta < 10^{-7}$, confirming numerical equivalence of the two implementations.




Figure S2: Training loss decomposition for the classical FBPINN baseline on the anomaly benchmark: total loss, PDE residual, displacement-snapshot initial conditions (two components), seismogram misfit, and free-stress boundary condition, plotted versus training iteration. Inset: early-iteration behaviour showing rapid initial loss reduction.

Supplementary Tables

Table S1: Full hyperparameter variant sweep for the anomaly benchmark. $\mathcal{Q}(n_q^{\times n_L})$ denotes a parameterized quantum circuit with n_q qubits and n_L layers. Values marked “_” are identical to the `classical` baseline. Subdomains column shows $\phi(\cdot)$ for wavefield decomposition and $c(\cdot)$ for velocity decomposition.

Model	ϕ network $(x, z, t) \rightarrow \phi$	c network $(x, z) \rightarrow c$	LR	Subs	Overlap	λ_{PDE}	λ_{IC}	λ_{seis}	λ_{BC}	# Pars
<code>classical</code>	$\mathcal{N}: 32^{\times 2} \rightarrow 16^{\times 2}$	$\mathcal{N}: 20^{\times 5}$	10^{-4}	$\phi(5 \times 3 \times 5)$	0.35	0.1	1.0	1.0	0.1	151,836
<i>ϕ-network size variants</i>										
<code>phi16x2</code>	$\mathcal{N}: 16^{\times 2}$	—	—	—	—	—	—	—	—	25,836
<code>phi64x3</code>	$\mathcal{N}: 64^{\times 3}$	—	—	—	—	—	—	—	—	649,836
<i>ϕ subdomain decomposition variants</i>										
<code>sub3x2x3</code>	—	—	—	$\phi(3 \times 2 \times 3)$	—	—	—	—	—	37,779
<code>sub8x4x8</code>	—	—	—	$\phi(8 \times 4 \times 8)$	—	—	—	—	—	514,017
<i>c-network size variants</i>										
<code>c10x2</code>	—	$\mathcal{N}: 10^{\times 2}$	—	—	—	—	—	—	—	150,226
<code>c10x3</code>	—	$\mathcal{N}: 10^{\times 3}$	—	—	—	—	—	—	—	150,336
<code>c20x3</code>	—	$\mathcal{N}: 20^{\times 3}$	—	—	—	—	—	—	—	150,996
<code>c40x3</code>	—	$\mathcal{N}: 40^{\times 3}$	—	—	—	—	—	—	—	153,516
<code>c40x5</code>	—	$\mathcal{N}: 40^{\times 5}$	—	—	—	—	—	—	—	156,796
<code>c60x3</code>	—	$\mathcal{N}: 60^{\times 3}$	—	—	—	—	—	—	—	157,636
<i>Decomposed c-network variants</i>										
<code>c_dec5x3</code>	—	$\mathcal{N}: 10^{\times 2}$	—	$\phi(5 \times 3 \times 5)$ $c(5 \times 3)$	—	—	—	—	—	152,340
<code>c_dec8x4</code>	—	$\mathcal{N}: 10^{\times 2}$	—	$\phi(5 \times 3 \times 5)$ $c(8 \times 4)$	—	—	—	—	—	154,907
<code>c_dec5x3_20</code>	—	$\mathcal{N}: 20^{\times 2}$	—	$\phi(5 \times 3 \times 5)$ $c(5 \times 3)$	—	—	—	—	—	156,600
<i>Training hyperparameter variants</i>										
<code>lr5e4</code>	—	—	5×10^{-4}	—	—	—	—	—	—	—
<code>lr5e5</code>	—	—	5×10^{-5}	—	—	—	—	—	—	—
<code>overlap50</code>	—	—	—	—	0.50	—	—	—	—	—
<code>wpde01_wbc01</code>	—	—	—	—	—	0.01	—	—	0.01	—
<code>wpde1_wbc1</code>	—	—	—	—	—	1.0	—	—	1.0	—
<i>Quantum hybrid</i>										
<code>quantum</code>	$\mathcal{N}: 32^{\times 2}$ $\rightarrow \mathcal{Q}: 4^{\times 2}$	$\mathcal{N}: 20^{\times 3}$ $\rightarrow \mathcal{Q}: 4^{\times 2}$	—	—	—	—	—	—	—	101,812